



UNIVERSITY OF MORATUWA

Faculty of Engineering
Department of Electronic and Telecommunication Engineering

EN3150

Machine Learning for Communication Systems

Assignment 02

Learning from Data and Related Challenges
and Classification

Submitted by:

Rivindu Vinsara Kumarage

Index No: 220343B

September 7, 2025

Contents

1	Linear Regression	2
2	Logistic Regression	6
3	Logistic Regression: First/Second-Order Methods	14
4	Conclusion	25

1 Linear Regression

Question 1 - OLS Line Alignment Issue

Analysis of OLS Misalignment

The Ordinary Least Squares (OLS) fitted line shown in Figure 1 is not aligned with the majority of data points due to the presence of **outliers** in the dataset.

Key Points:

- OLS minimizes the sum of squared residuals: $\frac{1}{N} \sum_{i=1}^N (y_i - \hat{y}_i)^2$
- The squared term gives disproportionate weight to large deviations (outliers)
- Outliers have a leverage effect, pulling the regression line towards them
- This results in a line that may not represent the true underlying relationship for the majority of the data

Mathematical Explanation: Since OLS uses squared errors, points far from the true line contribute quadratically to the loss function. For example, if a point has an error of 10 units, it contributes 100 to the loss, while a point with error 1 contributes only 1. This makes the algorithm prioritize reducing large errors over small ones, causing the fitted line to be biased towards outliers.

Question 2 - Modified Loss Function Analysis

Scheme Comparison

Scheme 1 will provide a better fitted line for inliers compared to the original OLS.

Analysis of Both Schemes:

Scheme 1: Outlier weight = 0.01, Inlier weight = 1

- Effectively reduces the influence of outliers by 99%
- Modified loss: $\frac{1}{N} \sum_{i=1}^N a_i (y_i - \hat{y}_i)^2$
- For outliers: $a_i = 0.01 \rightarrow$ minimal contribution to loss
- For inliers: $a_i = 1 \rightarrow$ normal contribution
- Result: Line fits primarily to inlier data

Scheme 2: Outlier weight = 5, Inlier weight = 1

- **Increases** the influence of outliers by 5×
- Makes the outlier problem worse than original OLS

- Line will be even more biased towards outliers
- Performs worse than standard OLS

Mathematical Justification: In Scheme 1, the optimization becomes:

$$\min_w \frac{1}{N} \left[\sum_{i \in \text{inliers}} (y_i - \hat{y}_i)^2 + 0.01 \sum_{i \in \text{outliers}} (y_i - \hat{y}_i)^2 \right]$$

This effectively ignores outliers and fits primarily to inlier data, resulting in better alignment with the majority of points.

Question 3 - Linear Regression Limitations in Brain Analysis

Why Linear Regression is Unsuitable

Linear regression is not suitable for identifying predictive brain regions due to several critical limitations:

1. High Dimensionality Problem

- Brain images contain thousands of voxels (features)
- Number of features \gg number of samples (curse of dimensionality)
- Standard linear regression cannot handle $p > n$ scenarios
- Results in overfitting and poor generalization

2. Multicollinearity Issues

- Adjacent voxels are highly correlated
- Creates unstable coefficient estimates
- Makes interpretation of individual voxel importance unreliable

3. Feature Selection Problem

- Cannot perform automatic feature selection
- All voxels receive non-zero weights
- Cannot identify specific brain regions as predictive
- Lacks sparsity in the solution

4. Interpretability Challenges

- Cannot group voxels into meaningful brain regions
- Individual voxel weights may not reflect regional importance
- Lacks biological/anatomical interpretability

Statistical Consequences: The design matrix $X \in \mathbb{R}^{N \times p}$ with $p \gg N$ makes $X^T X$ singular, preventing computation of $(X^T X)^{-1}$ needed for the closed-form solution.

Question 4 & 5 - LASSO vs Group LASSO Comparison

Method Selection and Justification

Group LASSO (Method B) is more appropriate for brain region identification.

Detailed Comparison:

Standard LASSO (Method A):

- Objective: $\min_w \left\{ \frac{1}{N} \sum_{i=1}^N (y_i - w^T x_i)^2 + \lambda \|w\|_1 \right\}$
- Performs individual voxel selection
- May select scattered voxels across brain regions
- Lacks spatial/anatomical coherence
- Cannot identify entire brain regions as predictive units

Group LASSO (Method B):

- Objective: $\min_w \left\{ \frac{1}{N} \sum_{i=1}^N (y_i - w^T x_i)^2 + \lambda \sum_{g=1}^G \|w_g\|_2 \right\}$
- Performs region-level selection (groups of voxels)
- Either selects entire brain regions or excludes them completely
- Maintains spatial/anatomical coherence
- Directly addresses the research question

Why Group LASSO is Superior:

1. Biological Relevance

- Brain functions are localized to specific regions
- Entire regions, not individual voxels, are functionally meaningful
- Group LASSO respects anatomical structure

2. Interpretability

- Directly identifies which brain regions are predictive
- Results are clinically/scientifically interpretable
- Facilitates hypothesis generation for neuroscience research

3. Statistical Advantages

- Reduces effective dimensionality from voxels to regions
- More stable feature selection
- Better handling of within-region correlation
- Improved prediction performance on new subjects

Mathematical Insight: The ℓ_2 norm penalty $\|w_g\|_2$ in Group LASSO creates an “all-or-nothing” selection mechanism for each group g . When λ is sufficiently large, entire groups are set to zero, effectively performing region-level feature selection that aligns with the research objective.

2 Logistic Regression

Question 1 - Data Loading

Data Loading

The provided code loads the penguins dataset and preprocesses it for binary classification between 'Adelie' and 'Chinstrap' species.

Code Analysis:

```
1 import seaborn as sns
2 import pandas as pd
3 from sklearn.model_selection import train_test_split
4 from sklearn.preprocessing import LabelEncoder
5 from sklearn.linear_model import LogisticRegression
6 from sklearn.metrics import accuracy_score
7
8 # Load and preprocess data
9 df = sns.load_dataset("penguins")
10 df.dropna(inplace=True)
11
12 # Filter for binary classification
13 selected_classes = ['Adelie', 'Chinstrap']
14 df_filtered = df[df['species'].isin(selected_classes)].copy()
15
16 # Encode labels
17 le = LabelEncoder()
18 y_encoded = le.fit_transform(df_filtered['species'])
19 df_filtered['class_encoded'] = y_encoded
```

Listing 1: Data Loading and Preprocessing

Question 2 - Training Errors with SAGA Solver

SAGA Solver Issues

Expected Errors:

1. Convergence Warning

- `ConvergenceWarning: lbfgs failed to converge`
- SAGA solver struggles with non-numerical features
- Default `max_iter` may be insufficient

2. Data Type Errors

- Categorical features ('island', 'sex') cause issues
- SAGA expects numerical input only
- String columns cannot be processed directly

Resolution Strategies:**Immediate Fixes:**

```
1 # Remove categorical columns
2 import numpy as np
3 X_numeric = df_filtered.select_dtypes(include=[np.number])
4 X_numeric = X_numeric.drop(['class_encoded'], axis=1)
5
6 # Or encode categorical variables
7 from sklearn.preprocessing import LabelEncoder
8 le_island = LabelEncoder()
9 le_sex = LabelEncoder()
10 df_filtered['island_encoded'] = le_island.fit_transform(
    df_filtered['island'])
11 df_filtered['sex_encoded'] = le_sex.fit_transform(df_filtered['
    sex'])
12
13 # Increase iterations
14 logreg = LogisticRegression(solver='saga', max_iter=1000)
```

Listing 2: Error Resolution

Comprehensive Preprocessing:

```
1 # Select only numerical features
2 numerical_features = ['bill_length_mm', 'bill_depth_mm',
3                       'flipper_length_mm', 'body_mass_g']
4 X = df_filtered[numerical_features]
```

Listing 3: Complete Preprocessing Pipeline

Question 3 - SAGA Solver Performance Issues

Why SAGA Performs Poorly**Technical Reasons for Poor Performance:****1. Algorithm Characteristics**

- SAGA is a stochastic gradient method
- Designed for very large datasets ($n > 10,000$)

- Inefficient for small datasets like penguins (~300 samples)
- High variance in gradient estimates for small samples

2. Feature Scaling Sensitivity

- SAGA is sensitive to feature scales
- Penguin features have different scales (mm vs grams)
- Unscaled features cause slow/poor convergence
- Gradient updates become imbalanced

3. Convergence Properties

- Requires many iterations to converge
- Default max_iter=100 often insufficient
- Stochastic nature leads to oscillating behavior
- Poor conditioning of the optimization landscape

Mathematical Explanation: SAGA updates follow: $w^{(k+1)} = w^{(k)} - \gamma(\nabla f_i(w^{(k)}) - \alpha_i^{(k)} + \bar{\alpha}^{(k)})$

For small datasets, the stochastic approximation adds unnecessary noise without computational benefits.

Question 4 - LibLinear Solver Performance

LibLinear Accuracy

Expected Classification Accuracy:

With the liblinear solver using only numerical features:

```

1 # Using numerical features only
2 X = df_filtered[['bill_length_mm', 'bill_depth_mm',
3                 'flipper_length_mm', 'body_mass_g']]
4 logreg = LogisticRegression(solver='liblinear')
5 logreg.fit(X_train, y_train)
6 accuracy = accuracy_score(y_test, y_pred)

```

Listing 4: LibLinear Implementation

Expected Accuracy: ~85-95%

The high accuracy is expected because:

- Adelie and Chinstrap penguins have distinct physical characteristics

- Clear separation in feature space
- Logistic regression is well-suited for this binary classification

Question 5 - LibLinear vs SAGA Comparison

Why LibLinear Outperforms SAGA

Algorithmic Advantages of LibLinear:

1. Optimization Method

- Uses coordinate descent algorithm
- Deterministic updates (no stochastic noise)
- Faster convergence for small-medium datasets
- More stable gradient estimates

2. Dataset Size Optimization

- Specifically designed for smaller datasets
- Efficient memory usage
- No overhead from stochastic sampling
- Better suited for $n < 10,000$ samples

3. Numerical Stability

- More robust to feature scaling issues
- Better conditioned optimization problem
- Consistent convergence behavior
- Less sensitive to hyperparameter choices

4. Implementation Efficiency

- Optimized C++ implementation
- Better cache locality
- Lower computational overhead per iteration
- Faster wall-clock time to convergence

Performance Summary:

Aspect	LibLinear	SAGA
Small datasets	Excellent	Poor
Convergence speed	Fast	Slow
Stability	High	Variable
Scaling sensitivity	Low	High

Question 6 - Random State Variance

Random State Impact on SAGA Accuracy

Sources of Variability:

1. Train-Test Split Randomness

- Different random_state values create different train/test splits
- Some splits may be more/less representative
- Creates baseline variability in performance measurement

2. SAGA Algorithm Stochasticity

- SAGA uses random sampling of gradients
- Different initialization leads to different optimization paths
- Convergence to different local optima possible
- High variance in final weight estimates

3. Convergence Issues

- May not converge within max_iter limit
- Stopping at different iteration counts
- Inconsistent solution quality
- Premature termination effects

Mathematical Explanation: The SAGA update rule introduces stochasticity:

$$w^{(k+1)} = w^{(k)} - \gamma \cdot [\nabla f_{i_k}(w^{(k)}) - \alpha_{i_k}^{(k)} + \bar{\alpha}^{(k)}]$$

Where i_k is randomly sampled, creating path-dependent convergence behavior.

Mitigation Strategies:

- Use cross-validation for robust evaluation

- Increase max_iter for better convergence
- Apply feature scaling
- Consider ensemble methods

Question 7 - Feature Scaling Impact

Scaling Comparison Analysis

Implementation and Results:

```

1 from sklearn.preprocessing import StandardScaler
2
3 # Without scaling
4 logreg_saga = LogisticRegression(solver='saga', max_iter=1000)
5 logreg_liblinear = LogisticRegression(solver='liblinear')
6
7 # With scaling
8 scaler = StandardScaler()
9 X_train_scaled = scaler.fit_transform(X_train)
10 X_test_scaled = scaler.transform(X_test)
11
12 logreg_saga_scaled = LogisticRegression(solver='saga', max_iter=1000)
13 logreg_liblinear_scaled = LogisticRegression(solver='liblinear')

```

Listing 5: Feature Scaling Comparison

Expected Results:

Solver	Without Scaling	With Scaling	Improvement
SAGA	60-75%	85-95%	Significant
LibLinear	85-90%	90-95%	Moderate

Reasons for Improvement:

1. SAGA Scaling Sensitivity

- Features have vastly different scales (mm vs grams)
- Gradient components become imbalanced
- Large-scale features dominate updates
- Scaling equalizes gradient contributions

2. Optimization Landscape

- Unscaled: elongated, ill-conditioned ellipsoids
- Scaled: more circular, well-conditioned contours
- Faster convergence along all dimensions
- More stable gradient descent

3. Differential Impact

- SAGA shows dramatic improvement (most sensitive)
- LibLinear shows modest improvement (more robust)
- Confirms algorithmic characteristics

Question 8 - Categorical Feature Encoding

Proper Categorical Feature Handling

The Approach is INCORRECT.

Problems with Label Encoding + Scaling:

1. Artificial Ordinality

- Label encoding: red=0, blue=1, green=2
- Implies: red < blue < green (false ordering)
- Creates meaningless distance relationships
- Scaling preserves these artificial relationships

2. Statistical Issues

- Standard scaling assumes continuous, normally distributed data
- Categorical labels are discrete and nominal
- Mean and standard deviation are meaningless
- Scaled values misrepresent categories

3. Model Interpretation Problems

- Logistic regression assumes linear relationships
- Coefficients become uninterpretable
- May learn spurious patterns from artificial ordering

Proposed Solutions:**Method 1: One-Hot Encoding (Recommended)**

```
1 from sklearn.preprocessing import OneHotEncoder
2 import numpy as np
3
4 # One-hot encoding for categorical features
5 encoder = OneHotEncoder(drop='first', sparse=False)
6 categorical_encoded = encoder.fit_transform(df[['color']].values
7     .reshape(-1, 1))
8
9 # Create feature names
10 feature_names = ['color_green', 'color_red'] # blue is dropped
11     (reference)
12
13 # Scale only numerical features
14 scaler = StandardScaler()
15 numerical_scaled = scaler.fit_transform(numerical_features)
16
17 # Combine features
18 X_final = np.hstack([numerical_scaled, categorical_encoded])
```

Listing 6: One-Hot Encoding Approach

Method 2: Mixed Preprocessing Pipeline

```
1 from sklearn.compose import ColumnTransformer
2 from sklearn.preprocessing import StandardScaler, OneHotEncoder
3
4 preprocessor = ColumnTransformer(
5     transformers=[
6         ('num', StandardScaler(), numerical_columns),
7         ('cat', OneHotEncoder(drop='first'), categorical_columns)
8     ]
9 )
10
11 X_processed = preprocessor.fit_transform(X)
```

Listing 7: Column Transformer Approach

Why This Approach is Superior:

- Preserves categorical nature of features
- No artificial ordinality introduced
- Proper statistical treatment for each feature type
- Maintains interpretability
- Standard scaling applied only where appropriate

3 Logistic Regression: First/Second-Order Methods

Question 1 - Data Generation

Data Generation Implementation

The provided code generates synthetic binary classification data using scikit-learn's `make_blobs` function with specific transformations.

```
1 import numpy as np
2 import matplotlib.pyplot as plt
3 from sklearn.datasets import make_blobs
4
5 np.random.seed(0)
6 centers = [[-5, 0], [5, 1.5]]
7 X, y = make_blobs(n_samples=2000, centers=centers, random_state
8                   =5)
9 transformation = [[0.5, 0.5], [-0.5, 1.5]]
10 X = np.dot(X, transformation)
```

Listing 8: Data Generation Analysis

This creates a linearly separable binary classification problem with transformed feature space.

Question 2 - Batch Gradient Descent Implementation

Gradient Descent Implementation

Weight Initialization Method: Xavier/Glorot Initialization is recommended for logistic regression.

Reasoning for Xavier Initialization:

- Maintains gradient magnitudes across layers
- Prevents vanishing/exploding gradient problems
- Accounts for number of input features
- Formula: $w \sim \mathcal{N}(0, \frac{1}{n_{\text{features}}})$

Implementation:

```
1 import numpy as np
2
3 def sigmoid(z):
4     """Sigmoid activation function with numerical stability"""
5     z = np.clip(z, -250, 250) # Prevent overflow
```

```
6     return 1 / (1 + np.exp(-z))
7
8 def batch_gradient_descent(X, y, n_iterations=20, learning_rate
   =0.01):
9     """
10    Batch Gradient Descent for Logistic Regression
11    """
12    # Add bias term
13    X_bias = np.c_[np.ones((X.shape[0], 1)), X]
14    n_samples, n_features = X_bias.shape
15
16    # Xavier initialization
17    np.random.seed(42)
18    weights = np.random.normal(0, 1/np.sqrt(n_features),
   n_features)
19
20    # Store loss history
21    loss_history = []
22
23    for i in range(n_iterations):
24        # Forward pass
25        z = X_bias.dot(weights)
26        predictions = sigmoid(z)
27
28        # Compute loss (cross-entropy)
29        epsilon = 1e-15 # Prevent log(0)
30        predictions = np.clip(predictions, epsilon, 1 - epsilon)
31        loss = -np.mean(y * np.log(predictions) +
32                        (1 - y) * np.log(1 - predictions))
33        loss_history.append(loss)
34
35        # Compute gradients
36        dw = (1/n_samples) * X_bias.T.dot(predictions - y)
37
38        # Update weights
39        weights -= learning_rate * dw
40
41        if i % 5 == 0:
42            print(f"Iteration {i}: Loss = {loss:.6f}")
43
44    return weights, loss_history
45
46 # Execute gradient descent
47 weights_gd, loss_gd = batch_gradient_descent(X, y, n_iterations
   =20)
```

Listing 9: Batch Gradient Descent Implementation

Alternative Initialization Methods:

- **Zero Initialization:** Simple but can cause slow convergence
- **Random Normal:** $w \sim \mathcal{N}(0, 0.01)$ - simple but less principled
- **He Initialization:** Better for ReLU, but Xavier is optimal for sigmoid

Question 3 - Loss Function Selection**Loss Function Choice****Selected Loss Function: Cross-Entropy (Log-Likelihood)**

Mathematical Form: $\mathcal{L}(w) = -\frac{1}{N} \sum_{i=1}^N [y_i \log(p_i) + (1 - y_i) \log(1 - p_i)]$

Where $p_i = \sigma(w^T x_i) = \frac{1}{1 + e^{-w^T x_i}}$

Reasons for Selection:**1. Statistical Foundation**

- Derived from maximum likelihood estimation
- Natural choice for probabilistic binary classification
- Provides proper probability estimates

2. Mathematical Properties

- Convex function (guaranteed global minimum)
- Smooth and differentiable everywhere and Well-behaved gradients for optimization

3. Gradient Characteristics

- Large gradients when predictions are wrong
- Small gradients when predictions are confident and correct
- Self-regulating learning behavior

4. Probabilistic Interpretation

- Outputs meaningful probabilities
- Enables uncertainty quantification
- Compatible with Bayesian inference

Question 4 - Newton's Method Implementation

Newton's Method Implementation

Implementation:

```

1 def newtons_method(X, y, n_iterations=20):
2     """
3     Newton's Method for Logistic Regression
4     """
5     # Add bias term
6     X_bias = np.c_[np.ones((X.shape[0], 1)), X]
7     n_samples, n_features = X_bias.shape
8
9     # Xavier initialization
10    np.random.seed(42)
11    weights = np.random.normal(0, 1/np.sqrt(n_features),
12                               n_features)
13
14    # Store loss history
15    loss_history = []
16
17    for i in range(n_iterations):
18        # Forward pass
19        z = X_bias.dot(weights)
20        predictions = sigmoid(z)
21
22        # Compute loss
23        epsilon = 1e-15
24        predictions_clipped = np.clip(predictions, epsilon, 1 -
25                                     epsilon)
26        loss = -np.mean(y * np.log(predictions_clipped) +
27                        (1 - y) * np.log(1 - predictions_clipped))
28        loss_history.append(loss)
29
30        # Compute first derivative (gradient)
31        gradient = (1/n_samples) * X_bias.T.dot(predictions - y)
32
33        # Compute second derivative (Hessian)
34        # H = (1/n) * X.T * W * X, where W is diagonal weight
35        # matrix
36        W = np.diag(predictions * (1 - predictions))
37        hessian = (1/n_samples) * X_bias.T.dot(W).dot(X_bias)
38
39        # Add regularization to prevent singular matrix
40        hessian += 1e-8 * np.eye(n_features)

```

```

39         try:
40             # Newton update: w = w - H-1 * gradient
41             delta_w = np.linalg.solve(hessian, gradient)
42             weights -= delta_w
43         except np.linalg.LinAlgError:
44             print(f"Singular matrix at iteration {i}, stopping")
45             break
46
47         if i % 5 == 0:
48             print(f"Iteration {i}: Loss = {loss:.6f}")
49
50     return weights, loss_history
51
52 # Execute Newton's method
53 weights_newton, loss_newton = newtons_method(X, y, n_iterations
54                                              =20)

```

Listing 10: Newton's Method for Logistic Regression

Key Implementation Details:**1. Hessian Computation**

- $H = \frac{1}{n} X^T W X$ where $W_{ii} = p_i(1 - p_i)$
- Represents curvature of loss function
- Always positive semi-definite for logistic regression

2. Numerical Stability

- Regularization term added to Hessian
- Using `np.linalg.solve()` instead of explicit inverse
- Clipping predictions to prevent numerical issues

3. Update Rule

- $w^{(k+1)} = w^{(k)} - H^{-1} \nabla \mathcal{L}(w^{(k)})$
- Uses second-order information for faster convergence
- Naturally adaptive step size

Question 5 - Convergence Comparison

Loss Comparison and Analysis

Plotting Implementation:

```
1 import matplotlib.pyplot as plt
2
3 plt.figure(figsize=(12, 8))
4 plt.plot(range(len(loss_gd)), loss_gd, 'b-o', linewidth=2,
5          markersize=6, label='Batch_Gradient_Descent', alpha
6          =0.8)
7 plt.plot(range(len(loss_newton)), loss_newton, 'r-s', linewidth
8          =2,
9          markersize=6, label="Newton's_Method", alpha=0.8)
10
11 plt.xlabel('Iteration', fontsize=14)
12 plt.ylabel('Cross-Entropy_Loss', fontsize=14)
13 plt.title('Convergence_Comparison:_Gradient_Descent_vs_Newton\'s
14           _Method',
15           fontsize=16, fontweight='bold')
16 plt.legend(fontsize=12)
17 plt.grid(True, alpha=0.3)
18 plt.yscale('log') # Log scale to better show convergence
19 plt.tight_layout()
20 plt.show()
21
22 # Print final losses
23 print(f"Final_GD_Loss:_{loss_gd[-1]:.8f}")
24 print(f"Final_Newton_Loss:_{loss_newton[-1]:.8f}")
25 print(f"GD_Iterations_for_convergence:_{len(loss_gd)}")
26 print(f"Newton_Iterations_for_convergence:_{len(loss_newton)}")
```

Listing 11: Loss Comparison Plot

Expected Results and Analysis:

1. Convergence Speed

- **Newton's Method:** Quadratic convergence (very fast)
- **Gradient Descent:** Linear convergence (slower)
- Newton typically converges in 3-8 iterations
- GD may require 50+ iterations for same precision

2. Loss Reduction Pattern

- **Newton:** Rapid exponential decrease initially
- **GD:** Steady linear decrease on log scale

- Newton shows steeper descent in early iterations
- Both eventually reach similar minimum

3. Mathematical Explanation

- Newton uses curvature information (Hessian)
- Better approximation of optimal step size
- GD uses fixed or simple adaptive learning rates
- Newton naturally handles ill-conditioning better

4. Computational Trade-offs

- **Newton:** $O(p^3)$ per iteration (Hessian inversion)
- **GD:** $O(p^2)$ per iteration (gradient computation)
- Newton faster for small-medium dimensions
- GD preferred for very high dimensions

Theoretical Convergence Rates:

- **Newton:** $\|w^{(k+1)} - w^*\| \leq C\|w^{(k)} - w^*\|^2$ (quadratic)
- **GD:** $\|w^{(k+1)} - w^*\| \leq \rho\|w^{(k)} - w^*\|$ where $\rho < 1$ (linear)

Question 6 - Stopping Criteria

Iteration Decision Approaches

Proposed Approaches for Stopping Criteria:

Approach 1: Convergence-Based Stopping

```

1 def convergence_based_stopping(loss_history, tolerance=1e-6):
2     """
3     Stop when loss change is below threshold
4     """
5     if len(loss_history) < 2:
6         return False
7
8     loss_change = abs(loss_history[-2] - loss_history[-1])
9     relative_change = loss_change / abs(loss_history[-2])
10
11     return relative_change < tolerance

```

```

12
13 # Usage in training loop
14 for i in range(max_iterations):
15     # ... training step ...
16     loss_history.append(current_loss)
17
18     if convergence_based_stopping(loss_history):
19         print(f"Converged at iteration {i}")
20         break

```

Listing 12: Convergence-Based Stopping

Advantages:

- Automatically adapts to convergence rate
- Prevents unnecessary computation
- Works for both GD and Newton's method
- Objective and reproducible

Approach 2: Validation-Based Early Stopping

```

1 def validation_based_stopping(X_train, y_train, X_val, y_val,
2                               patience=5, min_delta=1e-4):
3     """
4     Stop when validation loss stops improving
5     """
6     best_val_loss = float('inf')
7     patience_counter = 0
8
9     for i in range(max_iterations):
10         # Train for one iteration
11         weights = train_one_iteration(X_train, y_train, weights)
12
13         # Evaluate on validation set
14         val_loss = compute_loss(X_val, y_val, weights)
15
16         if val_loss < best_val_loss - min_delta:
17             best_val_loss = val_loss
18             patience_counter = 0
19         else:
20             patience_counter += 1
21
22         if patience_counter >= patience:
23             print(f"Early stopping at iteration {i}")
24             break
25

```

```
26 return weights
```

Listing 13: Validation-Based Early Stopping

Advantages:

- Prevents overfitting
- Better generalization performance
- Robust to noise in training data
- Standard practice in machine learning

Comparison of Approaches:

Criteria	Convergence-Based	Validation-Based
Prevents Overfitting	No	Yes
Computational Efficiency	High	Medium
Generalization	Unknown	Better
Implementation Complexity	Simple	Moderate
Data Requirements	Training only	Train + Validation

Recommendations:

- Use convergence-based for well-conditioned problems
- Use validation-based for real-world applications
- Combine both: convergence as backup, validation as primary
- Consider gradient norm: $\|\nabla \mathcal{L}\| < \epsilon$

Question 7 - Modified Centers Analysis**Convergence Analysis with New Centers****Modified Data Generation:**

```

1 # New centers configuration
2 centers_new = [[2, 2], [5, 1.5]]
3 X_new, y_new = make_blobs(n_samples=2000, centers=centers_new,
4                           random_state=5)
5 X_new = np.dot(X_new, transformation)
6
7 # Apply gradient descent
8 weights_new, loss_new = batch_gradient_descent(X_new, y_new,
```

```
n_iterations=20)
```

Listing 14: Modified Centers Configuration

Convergence Behavior Analysis:

1. Geometric Changes

- **Original centers:** $[-5, 0]$, $[5, 1.5]$ - well separated
- **New centers:** $[2, 2]$, $[5, 1.5]$ - closer together
- Reduced inter-class distance
- Increased class overlap after transformation

2. Impact on Separability

- Classes become less linearly separable
- Decision boundary becomes less obvious
- Higher inherent classification difficulty
- More complex optimization landscape

3. Expected Convergence Behavior

Slower Convergence:

- Gradients become smaller near decision boundary
- More iterations required for same precision
- Loss plateaus at higher minimum value
- Less steep descent in loss function

4. Mathematical Explanation

Gradient Magnitude Analysis: For logistic regression, gradient magnitude depends on prediction confidence: $\|\nabla \mathcal{L}\| = \frac{1}{n} \|X^T(p - y)\|$

When classes overlap:

- Predictions p closer to 0.5 (uncertain)
- Smaller values of $|p - y|$ for misclassified points
- Reduced gradient magnitudes
- Slower parameter updates

5. Visualization and Analysis


```

1 # Compare convergence rates
2 plt.figure(figsize=(15, 5))
3
4 # Plot 1: Data visualization
5 plt.subplot(1, 3, 1)
6 plt.scatter(X[:, 0], X[:, 1], c=y, alpha=0.7, cmap='viridis')
7 plt.title('Original Data (Well Separated)')
8 plt.xlabel('Feature 1')
9 plt.ylabel('Feature 2')
10
11 plt.subplot(1, 3, 2)
12 plt.scatter(X_new[:, 0], X_new[:, 1], c=y_new, alpha=0.7, cmap='
    viridis')
13 plt.title('Modified Data (Closer Centers)')
14 plt.xlabel('Feature 1')
15 plt.ylabel('Feature 2')
16
17 # Plot 3: Loss comparison
18 plt.subplot(1, 3, 3)
19 plt.plot(loss_gd, 'b-', label='Original Centers', linewidth=2)
20 plt.plot(loss_new, 'r--', label='Modified Centers', linewidth=2)
21 plt.xlabel('Iteration')
22 plt.ylabel('Loss')
23 plt.title('Convergence Comparison')
24 plt.legend()
25 plt.yscale('log')
26
27 plt.tight_layout()
28 plt.show()

```

Listing 15: Convergence Comparison Analysis

Expected Results:

Metric	Original	Modified
Final Loss	~ 0.01	$\sim 0.3 - 0.5$
Convergence Rate	Fast	Slow
Iterations to Converge	~ 10	$\sim 20+$
Classification Accuracy	$> 95\%$	$75 - 85\%$

6. Theoretical Implications**Condition Number Impact:**

- Closer centers increase condition number of Hessian
- Worse-conditioned optimization problem

- Requires smaller learning rates for stability
- Benefits more from second-order methods (Newton)

Practical Recommendations:

- Use adaptive learning rates (AdaGrad, Adam)
- Consider feature engineering for better separation
- Apply regularization to prevent overfitting
- Use early stopping based on validation performance

4 Conclusion

This assignment demonstrates the fundamental principles of machine learning optimization and classification techniques. Through theoretical analysis and practical implementation, we explored:

Key Learning Outcomes:

- Understanding of outlier impacts on regression methods
- Comparison of regularization techniques (LASSO vs Group LASSO)
- Practical solver selection for logistic regression
- Implementation of first and second-order optimization methods
- Analysis of convergence behavior under different data conditions

The results highlight the importance of proper preprocessing, algorithm selection, and understanding the underlying mathematical principles for successful machine learning applications.

References

- [1] Robert Tibshirani, “Regression shrinkage and selection via the lasso,” *Journal of the Royal Statistical Society Series B: Statistical Methodology*, vol. 58, no. 1, pp. 267–288, 1996.
- [2] Lukas Meier, Sara Van De Geer, and Peter Bühlmann, “The group lasso for logistic regression,” *Journal of the Royal Statistical Society Series B: Statistical Methodology*, vol. 70, no. 1, pp. 53–71, 2008.
- [3] Ming Yuan and Yi Lin, “Model selection and estimation in regression with grouped variables,” *Journal of the Royal Statistical Society Series B: Statistical Methodology*, vol. 68, no. 1, pp. 49–67, 2006.